

Manufacturing Consent for Israel’s Aggression: A Computational Analysis of Anti-Palestinian Bias in English-Language Media Headlines

Francisca Anguita^{*1}, Noah Allstedt^{*†1}, and Lenn Muluken^{*†1}

¹Social Science Research Master, University of Amsterdam, Roetersstraat 11,
Amsterdam, 1018 WB, The Netherlands

***Corresponding author(s):** francisca.anguita.magnis@student.uva.nl;
noah.allstedt@student.uva.nl; lenn.muluken@student.uva.nl

[†]These authors contributed equally to this work.

Abstract

This study investigates anti-Palestinian bias in Western media coverage of the Israel-Palestine conflict, focusing on how linguistic patterns in headlines shape public understanding and reproduce dominant power relations. Grounded in Critical Discourse Analysis and Orientalism, we analyze 1,624 headlines from The New York Times, BBC, and The Guardian between October 2023 and April 2025 using LLM and NLP techniques. GPT-4 was used to identify victims, whether a perpetrator is named, and whether the headline is written in an active or passive voice. Results show that there is relative underreporting of Palestinian suffering, it is more likely described using passive sentence constructions, and Palestinians are rarely referred to as ‘Palestinian’. These discursive choices downplay Israeli responsibility and obscure the scale and nature of Palestinian suffering. We argue that these patterns reflect broader ideological structures embedded in Western discourse, reinforcing Orientalist narratives and contributing to the manufacturing of consent for Israeli state violence. Our findings highlight the critical role of media discourse in legitimizing political dominance.

Keywords: Media Bias, Palestine, Israel, Critical Discourse Analysis, Orientalism, Large Language Models

1 Introduction

As Israel continues its aggression on Gaza, a UN Special Committee concluded that Israel’s methods of warfare are consistent with the characteristics of genocide [1]. Yet this war is not only fought with weapons on the ground; it is also fought through competing narratives. In Western countries like the US and the UK, which maintain close ties to Israel [2], the media plays a powerful role in shaping how the war is understood. Within this context, the way violence is framed becomes a crucial aspect in shaping public opinion and ultimately, the extent of governmental support for Israel’s war on Gaza [3]. Given the media’s influential role in this conflict, this study seeks to examine how Western media differs in their coverage of Israeli and Palestinian suffering.

Positioning ourselves within the tradition of Critical Discourse Analysis (CDA), we view the media as a contested arena in which competing narratives struggle over the reproduction of power relations and the exercise of dominance [4, 5].

Therefore, we analyze differences in linguistic patterns between reports about Israeli and Palestinian suffering. Our analysis focuses on Western English-language news media - namely the BBC, The Guardian, and The New York Times - due to their high reputation and dominant role in shaping global discourses. As agenda-setters - i.e. entities that have access to and control over the discourse - these media outlets wield significant influence over how readers interpret who bears responsibility [6]. We analyze news headlines, the most widely consumed and therefore impactful component of a news report. Taiwo [7] and van Dijk [8] argue that headlines serve as condensed ideological signals; they frame the interpretation of the text that follows, anchor attention, and often circulate independently.

Although studies on media coverage of the Israel-Palestine conflict date back to the 1980s [9], scholarly interest in media bias grew during the First Intifada, with early research highlighting how Israeli media marginalized Palestinian voices and framed Israeli violence as defensive [10, 11]. Around 2000, focus shifted to international media, especially in the U.S., where studies by Chomsky [9] and Zelizer et al. [12] found consistent pro-Israel and anti-Palestinian framing. While much of the early literature identified anti-Palestinian bias, some research [13, 14] documented anti-Israel bias. The vast majority of the aforementioned studies employ distinctly qualitative methods to analyze media bias. While qualitative language analysis allows for theoretically grounded insights and yields results with high interpretive validity, it is inherently limited in the volume of data that can be analyzed. This limitation increases the risk of cherry-picking and reduces the generalizability of the findings to other settings [15]. Studies that employ computational methods to analyze large volumes of text remain the exception. Jackson [16] uses Natural Language Processing (NLP) techniques - such as active/passive voice detection and sentiment analysis - to examine over 33,000 New York Times articles from the First and Second Intifada, finding a persistent and intensifying anti-Palestinian bias. Garcia-Corral et al. [17] use large language models (LLMs) to examine the use of causal language during the current aggression, comparing BBC reporting on the Israel-Palestine war with coverage by other media outlets, as well as with reporting on the Russia-Ukraine war. Their analysis reveals that the BBC frequently avoided attributing responsibility for Palestinian suffering to Israeli actions. Similarly, Sadiya et al. [18]

employ LLMs to annotate news content for bias, finding a disproportionate bias against Palestine across a range of media sources.

While the existing studies using computational methods to examine Western news coverage of the current Israel-Palestine war offer initial evidence of an anti-Palestinian bias, we identified three limitations which this paper seeks to improve upon. First, they largely lack a clear theoretical foundation. While advanced computational techniques are used to detect patterns of bias, these analyses are frequently not grounded in theory. Second, the studies fail to adequately interpret their findings. As a result, they remain largely descriptive, overlooking the broader ideological structures and power dynamics that shape media coverage. Finally, these studies are limited in scope and complexity, as they only focus on a single manifestation of bias in news coverage. However, power dynamics are embedded in language in multiple and complex ways, necessitating a more comprehensive examination of the various forms such bias can take. While we do not claim to offer an exhaustive account of all possible manifestations of media bias, we provide a more nuanced analysis of such.

Our paper builds on CDA, particularly the work of Van Dijk [8, 19], who argues that the content and structure of texts represent a form of power through which discourse enables hegemony. Central to CDA is the idea that language actively shapes and is shaped by the world around us [4]. Within this framework, media outlets emerge as critical sites for the (re-)production of dominance. Media discourse thus becomes a linguistic tool, strategically shaped to serve the interests of those in power. Van Dijk [19] demonstrated how Western media uses strategies such as lexical choice and active/passive voice to dehumanize ethnic minorities and immigrants [20, 8, 21].

CDA views language as inherently non-neutral. Journalists, which are embedded within dominant discursive structures, make (un-)conscious linguistic choices in how they present events. For example, employing the active voice assigns agency, while passive constructions obscure responsibility. Furthermore, Taiwo [7] argues that terms such as “terrorist” do more than describe - they moralize. Moreover, it is not only a matter of how something is said, but whether it is said at all. That is, when the media reports a killing without naming the perpetrator, it is not simply omitting a detail; rather it is actively erasing accountability.

Our focus on the representation of Palestinian suffering is further grounded in Edward Said’s [22] theory of Orientalism. Said conceptualizes Orientalism as discourse that constructs the “Orient” not merely as a place but as a set of negative characteristics, such as violent and underdeveloped. By contrast, the West is constructed as rational and civilized. This binary serves to justify colonial domination and Western intervention. The Orient thus functions ideologically, it is applied to groups outside the West and helps reproduce hierarchies of power. Said expanded his critique in *Covering Islam* [23], where he shows how journalists frame events through stereotypes, portraying Arabs as threats. This notion is supported by Shaheen [24] who examined how Hollywood has stereotyped Arabs as villains and terrorists, reinforcing orientalist tropes. Further, Karim [25] and Kumar [26], showed how Western news discourse constructs the “Orient” as incompatible with modernity and Philo and Berry [27] demonstrated that Palestinian casualties are often decontextualized

and depersonalized.

Building on CDA and Said’s notion of Orientalism, this study examines linguistic patterns in Western media coverage to reveal the subtle yet powerful ways in which language influences how the war is understood.

2 Methodology

We use OpenAI’s GPT-4 and traditional NLP techniques to examine how Western media reports since the Hamas attack. Our dataset entails New York Times headlines that we obtained using their publicly available API and was supplemented with BBC and Guardian headline datasets that were publicly available. The BBC and Guardian headlines cover the period from October 7 2023 until the end of 2024, whereas the New York Times headlines run until April 30th 2025. Of 5961 headlines which report on the Israel-Palestine war, 1624 are kept for the analysis after filtering for a wordbank with words that describe various forms of suffering, such as death, injury or arrest (for full list, see Appendix A) .

We use GPT-4 to perform the following three tasks:

1. Identifying whether the described victim in the headline is Palestinian, Israeli or neither.
2. Identifying whether the perpetrator is explicitly mentioned and classifying who the perpetrator is.
3. Classifying whether the headline is written in active or passive voice (Active voice is when the subject performs the action, while passive voice is when the subject receives the action).

We further employ a keyword-matching approach to assess whether group labels ("Palestinian" or "Israeli") appear in headlines describing corresponding instances of suffering.

We initially attempted to employ traditional NLP techniques (spaCy and Stanza) to identify voice and the mentioning of a perpetrator, given that their rule-based nature ensures transparency and facilitates the replicability of findings. However, NLP techniques struggled with atypical sentence structures often found in media headlines, such as passive constructions that lack auxiliary verbs (e.g. 'Hundreds Reported Killed'). Hence, we choose to use GPT-4. This aligns with previous research demonstrating that LLMs outperform traditional NLP techniques, particularly in complex tasks that require contextual understanding [28, 29, 30]

We prompt GPT-4 separately for each task. To ensure the quality of the output, we utilize multiple strategies that yield the most accurate output. First, the LLM is prompted with a role, such as "You are a grammar analyst.". Prior research indicates that role-prompting significantly improves the accuracy of GPT-4 [31]. Second, for the task of identifying the victim, one-shot prompting is used. In contrast, few-shot prompting is employed for the tasks of identifying explicitly mentioned perpetrators and classifying headlines as using active or passive voice, in order to account for the higher complexity of the task and to better handle special cases [31]. This is particularly important for the voice classification task, as some headlines include

multiple clauses containing both active and passive constructions (e.g. “At least half of Gaza’s buildings are damaged or destroyed, new analysis shows.”). In such cases, GPT-4 is instructed to focus on the clause conveying the main claim, rather than the source attribution. (For full prompt examples, see Appendix B.) Finally, the LLM is given a detailed task description to ensure accurate and consistent outputs [32, 31]. The prompt engineering for the task specification follows an iterative process. For each task, the output of the LLM is compared to a random set of 50 manually annotated headlines. The prompt is then refined until a satisfactory level of agreement is achieved. To assess the level of agreement between the GPT-generated annotations and the manually coded sample for each individual task we calculate the inter-coder reliability, using Krippendorff’s alpha (see Appendix C).

Table 1: Inter-annotator agreement metrics for each variable.

Variable	Accuracy (%)	Krippendorff’s Alpha
Suffering	80.0	0.679
Perpetrator	88.0	0.821
Active/Passive	84.0	0.623

Two ethical concerns must be addressed. First, our analysis complies with GDPR regulations and addresses ethical concerns regarding data consent, as we rely solely on publicly available data that do not compromise individual privacy. Therefore, that OpenAI receives access to all data transmitted through its API does not pose a problem [32]. Second, a frequently raised concern regarding the use of LLMs is that they reflect the biases present in their training data [33, 32]. In addition to these data-driven biases, the personal biases of the engineers involved in building and fine-tuning the models may also influence the outputs. Given the contested nature of our topic, some degree of bias in our results is almost certain. This concern is supported by studies that have identified anti-Muslim bias in GPT models [34, 35]. Nonetheless we cannot know the form or direction in which these biases manifest, as proprietary LLMs like GPT-4 constitute black boxes.

3 Results

3.1 Reporting Suffering

Out of 1,624 news headlines analyzed, 653 headlines describe Palestinian suffering, followed and 516 headlines describe Israeli suffering. These results indicate a slightly greater prevalence of reporting on Palestinian suffering compared to Israeli suffering in the sample.

This apparent similarity in coverage becomes disproportionate when contextualized within the official death toll, where Palestinian casualties have been estimated at at least 53,655 whereas Israeli casualties are estimated to be at 1,616 [36]. Therefore, these results suggest that despite experiencing higher levels of violence and loss, Palestinian victims are only marginally more visible in media headlines.

Table 2: *Distribution of Suffering Labels in Sample*
($N = 1,624$)

Suffering Label	Frequency
Palestinian suffering	653
Israeli suffering	516
No suffering described	263
Other suffering	192

3.2 Reporting of Perpetrator and Victim

Figure 1 depicts how frequently headlines reporting on Palestinian and Israeli suffering explicitly identify the perpetrator. The difference in mentioning the perpetrator is marginal: perpetrators are mentioned in 51.9% of all headlines describing Palestinian suffering, compared to 47.7% for Israeli suffering. Figure 2, which illustrates the proportion of headlines that explicitly use the labels “Palestinian” and “Israeli” when describing suffering on each side, reveals a different picture. While Israeli victims are identified as such in 39% of the cases, Palestinian victims are only referred to as Palestinian in 15% of the headlines.




Figure 1 (Left) Proportion of headlines omitting perpetrators. **Figure 2** (Right) Proportion of headlines using group labels for victims.

3.3 Active and Passive Voice

Next, we plot the proportions of coverage that uses active/passive voice in suffering-related headlines. We create two plots, one that visualizes the voice proportions for all suffering-related headlines and a second that illustrates voice proportions on a subsample for death-related headlines only. As shown in Figure 3, there is a strong disparity in voice usage between reports on Israeli and Palestinian suffering. While 77% of headlines covering Israeli suffering use the active voice, only 59% do so in the case of Palestinian suffering. A similar, albeit, less pronounced pattern can be observed in the headlines that specifically mention death (Figure 4). While Israeli deaths are described in the active voice in 47% of the headlines, the active voice is only used in 40% of the headlines that report on Palestinian deaths.




Figure 3 (Left) Proportion of headlines in active/passive voice for all headlines related to suffering. **Figure 4** (Right) Proportion of headlines in active/passive voice for all headlines related to deaths

3.4 Reported Suffering in Headlines Over Time

Figure 5 plots the number of headlines that report on Palestinian and Israeli suffering over time, aggregated by week starting from October 2023 to May 2025. Both Palestinian and Israeli suffering received extensive coverage in the initial weeks, with Israeli suffering peaking after the Hamas attack on October 7th. Palestinian suffering headlines followed shortly after, as Palestinians deaths increased due to Israel’s attacks on Gaza.

One visible trend that appears when observing coverage over time is that coverage on Palestinian suffering and Israeli suffering follow similar trends. Spikes in coverage of suffering often converge. This appears somewhat surprising, considering that the Israeli death toll remained relatively unchanged after October 7th [36].




Figure 5: Media Headline count on Israeli and Palestinian Suffering for each month

4 Discussion

Interpreting our findings through the lens of CDA, we argue that media coverage of the Israel-Palestine conflict (1) constructs social reality through discourse and (2) consequently acts as an instrument to reproduce existing power relations.

One way this manifests is through the relative underreporting of Palestinian suffering in the media, which functions to convey a false balance of the war. Thus, Palestinian suffering is rendered less news-worthy which aligns with Orientalist tropes in which Arab lives are dehumanized and stripped of their individuality. The dehumanization of Arab lives in broader Western discourses suggests that Palestinian casualties are indistinguishable and do not require separate media attention [23]. Furthermore, by using passive language to describe Palestinian suffering Western media erases Israel’s responsibility for its aggression. When media headlines do report about Palestinian suffering, the frequent use of passive language obscures how said suffering came about. Therefore, passive structures fail to question Israeli actions. A further finding is that headlines describing Palestinian suffering are significantly less likely to explicitly label victims as ‘Palestinians’ compared to those describing Israeli suffering. Avoiding to label Palestinians as such discursively denies their identity. In the context of Israel’s expansionary politics and continuous moves to gradually annex the Palestinian territories, the non-existence of a Palestinian national identity erases the right of Palestinians to their land. By extension this coverage reinforces the Israeli narrative, according to which historically Palestinians do not exist in order to justify resettlement and ethnic cleansing [37].

Moreover, as discussed, spikes in coverage about Palestinian suffering frequently coincide with coverage of Israeli suffering. This potentially indicates that whenever Israel acts particularly aggressive towards Palestinians, Israeli suffering receives coverage to serve as a justification for Israeli aggressions. However, since our analysis does not directly address the causes of this pattern, our interpretation remains speculative. Further qualitative research is needed to provide a better understanding of the underlying drivers.

The different linguistic and coverage choices in the three media outlets relativize and neglect Palestinian suffering, mask Israeli responsibility and ultimately serve to support the Israeli narrative. Through the CDA lens, the observed discursive patterns are not merely accidental but are structural manifestations in the West, which views Muslims in the Middle East as violent and inferior and Israel as a Western-oriented ally. This has real-world material consequences, as media language can influence public opinion and consequently impact government policy. Therefore, media coverage contributes to the manufacturing of consent for policies harmful to Palestinians such as weapon exports and allows Israel to act without impunity.

Our study is not without limitations. First, the dataset contains some noise. Although we restricted our word bank to terms primarily associated with the Israel–Palestine conflict, some unrelated headlines were inadvertently included. Conversely, it is also possible that relevant headlines were omitted if they did not contain the selected keywords. However, as long as this occurs at random, they are unlikely to introduce significant bias into our findings. Second, while our methodological approach allowed us to identify broader patterns in media coverage, it could not provide a more nuanced analysis of individual headlines. Future qualitative research could delve deeper into the specific use of language to uncover idiosyncratic discursive strategies. Finally, while our use of CDA enables us to analyze underlying power structures we cannot offer certain evidence of their existence due to their latent nature. However, CDA explicitly urges its followers not to shy away from being bold scientists. CDA argues that there is no neutral social science and therefore chooses to be explicitly political.

5 Conclusion

The present study examined whether an anti-Palestinian bias exists across three high-quality Media outlets by comparing differences in reporting on Israeli and Palestinian suffering since October 7th, 2023. It analyzed the overall amount of coverage of suffering on each side, the mention of perpetrators and victims, as well as the use of active versus passive voice, and found evidence for an anti-Palestinian bias in all but one dimension (the mention of perpetrators). The main argument of this paper is that Western media outlets do not reinforce the Israeli narrative by mere chance, but that this occurs structurally within the West, which views Muslims in the Middle East as backwards and maintains close ties to the Israeli state.

6 Data, Materials, and Software Availability

Appendices, the full Python code, and the datasets have been deposited in GitHub, and is available from: <https://github.com/LennMuluken/CSSProject.git>

7 Acknowledgement

We thank the anonymous peer-reviewers for their helpful and constructive feedback. ChatGPT was used to aid the coding process in Python. It was employed to help extract NYT headlines with the NYT API and to build three GPT-4 pipelines for automated classifications using the ChatCompletion interface.

References

- [1] OHCHR. *UN Special Committee Finds Israel's Warfare Methods in Gaza Consistent with Genocide*. <https://www.ohchr.org/en/press-releases/2024/11/un-special-committee-finds-israels-warfare-methods-gaza-consistent-genocide>. 2024.
- [2] CNN. *Where does Israel get its weapons?* Accessed: 2025-05-28. 2024. URL: <https://edition.cnn.com/2024/10/16/middleeast/where-israel-get-its-weapons>.
- [3] Herman, Edward S. and Chomsky, Noam. *Manufacturing Consent*. Routledge, 2021, pp. 198–206.
- [4] Fairclough, Norman. “Critical Discourse Analysis”. In: *The Routledge Handbook of Discourse Analysis*. 1st. London: Routledge, 2013, pp. 9–20. DOI: 10.4324/9780203809068.
- [5] Dijk, Teun A. van. *Critical Discourse Analysis*. <https://discourses.org/wp-content/uploads/2022/07/Teun-A.-van-Dijk-2015-Critical-discourse-Analysis.pdf>. 2015.
- [6] Wanta, Wayne and Hu, Yu-Wei. “The effects of credibility, reliance, and exposure on media agenda-setting: A path analysis model”. In: *Journalism Quarterly* 71.1 (1994), pp. 90–98.

- [7] Taiwo, Rotimi. “Language, Ideology and Power Relations in Nigerian Newspaper Headlines”. In: *Nebula* 4.1 (2007), pp. 218–245.
- [8] Dijk, Teun A. van. *Racism and the Press*. Routledge, 1991.
- [9] Chomsky, Noam. *Fateful Triangle: The United States, Israel, and the Palestinians*. South End Press, 1999.
- [10] Wolfsfeld, Gadi. *Media and Political Conflict: News from the Middle East*. Cambridge University Press, 1997.
- [11] Liebes, Tamar. *Reporting the Arab-Israeli Conflict*. 1997.
- [12] Zelizer, Barbie, Park, David, and Gudelunas, David. “How bias shapes the news: Challenging The New York Times’ status as a newspaper of record on the Middle East”. In: *Journalism* 3.3 (2002), pp. 283–307.
- [13] Segev, Elad and Miesch, R. “A systematic procedure for detecting news biases: The case of Israel in European news sites”. In: *International Journal of Communication* 5 (2011), p. 20.
- [14] Neureiter, Michael. “Sources of media bias in coverage of the Israeli–Palestinian conflict: the 2010 Gaza flotilla raid in German, British, and US newspapers”. In: *Israel Affairs* 23.1 (2017), pp. 66–86.
- [15] Firestone, William A. “Alternative Arguments for Generalizing from Data as Applied to Qualitative Research”. In: *Educational Researcher* 22.4 (1993), pp. 16–23.
- [16] Jackson, Hannah M. “The New York Times Distorts the Palestinian Struggle: A Case Study of Anti-Palestinian Bias in US News Coverage of the First and Second Palestinian Intifadas”. In: *Media, War Conflict* 17.1 (2024), pp. 116–135.
- [17] Garcia-Corral, Paula et al. “The Missing Cause: An Analysis of Causal Attributions in Reporting on Palestine”. In: *The 31st International Conference on Computational Linguistics*. Jan. 2025, p. 103.
- [18] Sadijah, N., Al-Emadi, S., and Rahman, S. “Ceasefire at FIGNEWS 2024 Shared Task: Automated Detection and Annotation of Media Bias Using Large Language Models”. In: *Proceedings of The Second Arabic Natural Language Processing Conference*. Aug. 2024, pp. 590–600.
- [19] Dijk, Teun A. van. “War rhetoric of a little ally”. In: *Journal of Language and Politics* 4.1 (2005), pp. 65–91. DOI: 10.1075/jlp.4.1.04dij.
- [20] Dijk, Teun A. van. *Prejudice in Discourse*. 1984.
- [21] Dijk, Teun A. van. “Elite Discourse and Racism”. In: *Approaches to Discourse, Poetics and Psychiatry*. Vol. 4. 1993, p. 81.
- [22] Said, Edward W. *Orientalism*. Penguin Classics, 2003.
- [23] Said, Edward W. *Covering Islam: How the Media and the Experts Determine how We See the Rest of the World*. New York: Pantheon Books, 1997.
- [24] Shaheen, Jack G. “Reel Bad Arabs: How Hollywood Vilifies People”. In: *The Annals of the American Academy of Political and Social Science* 588 (2003). <http://www.jstor.org/stable/1049860>, pp. 171–193.

- [25] Karim, Karim H. *Islamic Peril: Media and Global Violence*. 2000.
- [26] Kumar, Deepa. *Islamophobia and the Politics of Empire*. Haymarket Books, 2012.
- [27] Philo, Greg et al. *Bad News From Israel*. <http://www.jstor.org/stable/j.ctt183q4ss>. Pluto Press, 2004.
- [28] Chang, Yuxin et al. “A Survey on Evaluation of Large Language Models”. In: *ACM Transactions on Intelligent Systems and Technology* 15.3 (2024), pp. 1–45.
- [29] Gilardi, Fabrizio, Alizadeh, Meysam, and Kubli, Marcel. “ChatGPT Outperforms Crowd Workers for Text-Annotation Tasks”. In: *Proceedings of the National Academy of Sciences* 120.30 (2023), e2305016120.
- [30] Lopez-Lira, Andrés and Tang, Youwei. *Can ChatGPT Forecast Stock Price Movements? Return Predictability and Large Language Models*. 2023. arXiv: 2304.07619 [q-fin.ST].
- [31] Chen, Bin et al. “Unleashing the Potential of Prompt Engineering for Large Language Models”. In: *Patterns* (2025).
- [32] Törnberg, Petter. “Best Practices for Text Annotation with Large Language Models”. In: *Sociologica (Bologna)* 18.2 (2023), pp. 67–85. DOI: 10.6092/issn.1971-8853/19461.
- [33] Gallegos, Ignacio O. et al. “Bias and Fairness in Large Language Models: A Survey”. In: *Computational Linguistics* 50.3 (2024), pp. 1097–1179.
- [34] Abid, Abubakar, Farooqi, Maheen, and Zou, James. “Persistent Anti-Muslim Bias in Large Language Models”. In: *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*. July 2021, pp. 298–306.
- [35] Hemmatian, Behnam, Baltaji, Rayan, and Varshney, Lav R. *Muslim-Violence Bias Persists in Debaised GPT Models*. 2023. arXiv: 2310.18368 [cs.CL].
- [36] OCHA. *Humanitarian Situation Update #290: Gaza Strip*. <https://www.ochaopt.org/content/humanitarian-situation-update-290-gaza-strip>. 2025.
- [37] Masalha, Nur. *Imperial Israel and the Palestinians: The Politics of Expansion*. Pluto Press, 2000.

8 Response To Peer Reviewers

Dear peer reviewers, thank you for your constructive feedback. The concerns voiced by the reviewers helped us to revise and improve our paper, especially our theory and discussion sections. For concerns that were not integrated we provide a justification below.

A major concern regarding our paper was the lack of theoretical rigour when discussing the behaviour of media outlets, journalists and societies, as certain claims were insufficiently substantiated with sources. We have since reviewed the paper carefully and integrated relevant literature to support our claims and rephrased statements to reflect a more cautious tone. Further, one reviewer observed that the discussion of Said’s Orientalism was underdeveloped. We have since clarified and

expanded our discussion of Said’s Orientalism. We also clarified how Orientalist discourse manifests in media representations by including appropriate sources which have extended Said’s framework to media discourse, thereby grounding our argument more substantially within scholarly debates.

Another concern was raised regarding the discussion and conclusion section, which included overreaching claims, particularly concerning the broader political implications of our findings. We have now revised and reframed this section to stay closer to the data. Statements were rephrased to reflect the discursive tendencies observed. In cases where conjectural arguments were made, we have either highlighted their tentativeness or strengthened the justification with clear links to the theory. We further highlight both dis-/advantages of CDA which encourages relating language to latent power structures.

Additionally, the peer reviewers disagreed with the interpretation of one graph depicting the coverage of Palestinian and Israeli suffering over time. We agree with the concern that our interpretation overstated the frequency and strength of the observed pattern and therefore revised that section so that our interpretation reflects the data more closely. While we still identify a trend, we now present our interpretation of said pattern as speculative. However, we maintain that this pattern, although tentative, is worth mentioning, as we consider it a valuable direction for future research to look into.

Moreover, we were advised to provide further justification for the choice of the media outlets. We would like to clarify that this decision is addressed in the introduction where we state that The New York Times, BBC, and The Guardian were selected due to their high reputation, large audience and dominant role in shaping international discourse. Further, we were limited to publicly available datasets and newspapers that provide access to their API.

Additionally, peer reviewers commented that our validation of the LLM results, based on a manually annotated sub-sample representing 1,8% of the total sample, may be insufficient to reliably assess the model’s accuracy. To address this concern we have expanded our annotated sample to 50 headlines for each prompt and conducted Krippendorff’s alpha tests to assess the reliability of GPT-4’s annotations.

Furthermore, our analysis was criticized for not employing a chi-square test to determine whether the observed differences are statistically significant. However, we do not consider such testing necessary in our case, as our dataset of headlines represents the entire population of interest rather than a representative sample drawn from it. Consequently, inferential statistical methods aimed at estimating uncertainty do not apply.

Moreover, we clarified that the NYT dataset runs until the end of April 2025, integrated that LLMs most definitely contain biases, and provided a definition of active/passive voice. We also briefly elaborated on the analysis used in papers that use computational methods to examine news coverage of the Israel-Palestine war.

Finally, we also addressed minor issues relating to layout, grammar, references, figure description, acronyms of central concepts, inclusion of appendices.